

```
In [5]: import os
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

## HW 5

```
In [6]: train = pd.read_csv("MusicRatingsTrain.csv")
test = pd.read_csv("MusicRatingsTest.csv")
vala = pd.read_csv("MusicRatingsValidationA.csv")
valb = pd.read_csv("MusicRatingsValidationB.csv")
train.head()
```

Out[6]:

|   | userID | songID | rating   | songName                    | year | artist              | genre      |
|---|--------|--------|----------|-----------------------------|------|---------------------|------------|
| 0 | 2071   | 581    | 2.484754 | Show You How                | 2004 | The Killers         | Rock       |
| 1 | 1208   | 55     | 1.000000 | Luvstruck                   | 1999 | Southside Spinners  | Electronic |
| 2 | 1484   | 510    | 1.000000 | In One Ear                  | 2008 | Cage The Elephant   | Rock       |
| 3 | 453    | 299    | 1.000000 | Gardenhead / Leave Me Alone | 1992 | Neutral Milk Hotel  | Rock       |
| 4 | 2238   | 113    | 1.000000 | Hustler                     | 2006 | Simian Mobile Disco | Electronic |

## Question 1

a

```
In [7]: songs = len(pd.unique(train['songID']))
songs
```

Out[7]: 807

```
In [8]: users = len(pd.unique(train['userID']))
users
```

Out[8]: 2421



```
In [9]: stats = train['rating'].describe()
stats
```

```
Out[9]: count      245997.000000
mean           1.321962
std            0.459894
min            1.000000
25%            1.000000
50%            1.000000
75%            1.494918
max            4.768656
Name: rating, dtype: float64
```

```
In [10]: median = train['rating'].median()
median
```

```
Out[10]: 1.0
```

## b) i.

There are a total of  $807 + 2421 = 3228$  parameters included in the model. We have to train the model with 245997 observations because there are no null ratings in the training dataset, so the number of observations is the same as the number of ratings.

## b) ii.

```
In [11]: from fancyimpute import BiScaler
```

```
In [12]: train.sort_index()
train_df = train.pivot_table(index="userID", columns = "songID", values = "
train_mat = train_df.to_numpy()
train_mat
```

```
Out[12]: array([[ nan,      nan,      nan, ...,      nan,      nan,
               [ nan, 1.98983574,      nan, ...,      nan,      nan,
               [ nan,      nan,      nan, ...,      nan,      nan,
               ...,
               [ nan,      nan,      nan, ...,      nan,      nan,
               [ nan,      nan,      nan, ...,      nan,      nan,
               [ nan, 1.49491787, 1.49491787, ...,      nan,      nan,
               nan]])
```

```
In [13]: song_biscaler = BiScaler(scale_rows=False, scale_columns=False, verbose=True)
train_mat_centered = song_biscaler.fit_transform(train_mat)
```

```
[BiScaler] Initial log residual value = 8.635876
[BiScaler] Iter 1: log residual = -1.278081, log improvement ratio=9.9139
58
[BiScaler] Iter 2: log residual = -2.464318, log improvement ratio=1.1862
37
[BiScaler] Iter 3: log residual = -3.599288, log improvement ratio=1.1349
70
[BiScaler] Iter 4: log residual = -4.728969, log improvement ratio=1.1296
81
[BiScaler] Iter 5: log residual = -5.851105, log improvement ratio=1.1221
36
[BiScaler] Iter 6: log residual = -6.961307, log improvement ratio=1.1102
02
[BiScaler] Iter 7: log residual = -8.052791, log improvement ratio=1.0914
84
[BiScaler] Iter 8: log residual = -9.115833, log improvement ratio=1.0630
43
[BiScaler] Iter 9: log residual = -10.137984, log improvement ratio=1.022
150
[BiScaler] Iter 10: log residual = -11.105953, log improvement ratio=0.96
7969
[BiScaler] Iter 11: log residual = -12.009465, log improvement ratio=0.90
3512
[BiScaler] Iter 12: log residual = -12.845445, log improvement ratio=0.83
5981
[BiScaler] Iter 13: log residual = -13.619498, log improvement ratio=0.77
4052
[BiScaler] Iter 14: log residual = -14.343386, log improvement ratio=0.72
3888
[BiScaler] Iter 15: log residual = -15.030646, log improvement ratio=0.68
7260
[BiScaler] Iter 16: log residual = -15.693214, log improvement ratio=0.66
2568
[BiScaler] Iter 17: log residual = -16.340088, log improvement ratio=0.64
6874
[BiScaler] Iter 18: log residual = -16.977415, log improvement ratio=0.63
7327
[BiScaler] Iter 19: log residual = -17.609141, log improvement ratio=0.63
1726
[BiScaler] Iter 20: log residual = -18.237697, log improvement ratio=0.62
8555
[BiScaler] Iter 21: log residual = -18.864535, log improvement ratio=0.62
6838
[BiScaler] Iter 22: log residual = -19.490503, log improvement ratio=0.62
5968
[BiScaler] Iter 23: log residual = -20.116082, log improvement ratio=0.62
5579
[BiScaler] Iter 24: log residual = -20.741536, log improvement ratio=0.62
5455
[BiScaler] Iter 25: log residual = -21.367004, log improvement ratio=0.62
5468
[BiScaler] Iter 26: log residual = -21.992551, log improvement ratio=0.62
5547
[BiScaler] Iter 27: log residual = -22.618202, log improvement ratio=0.62
```

5651  
[BiScaler] Iter 28: log residual = -23.243962, log improvement ratio=0.62  
5760  
[BiScaler] Iter 29: log residual = -23.869823, log improvement ratio=0.62  
5861  
[BiScaler] Iter 30: log residual = -24.495775, log improvement ratio=0.62  
5951  
[BiScaler] Iter 31: log residual = -25.121804, log improvement ratio=0.62  
6029  
[BiScaler] Iter 32: log residual = -25.747898, log improvement ratio=0.62  
6094  
[BiScaler] Iter 33: log residual = -26.374046, log improvement ratio=0.62  
6148  
[BiScaler] Iter 34: log residual = -27.000238, log improvement ratio=0.62  
6192  
[BiScaler] Iter 35: log residual = -27.626466, log improvement ratio=0.62  
6228  
[BiScaler] Iter 36: log residual = -28.252723, log improvement ratio=0.62  
6257  
[BiScaler] Iter 37: log residual = -28.879004, log improvement ratio=0.62  
6281  
[BiScaler] Iter 38: log residual = -29.505303, log improvement ratio=0.62  
6299  
[BiScaler] Iter 39: log residual = -30.131618, log improvement ratio=0.62  
6314  
[BiScaler] Iter 40: log residual = -30.757944, log improvement ratio=0.62  
6326  
[BiScaler] Iter 41: log residual = -31.384280, log improvement ratio=0.62  
6336  
[BiScaler] Iter 42: log residual = -32.010624, log improvement ratio=0.62  
6344  
[BiScaler] Iter 43: log residual = -32.636973, log improvement ratio=0.62  
6350  
[BiScaler] Iter 44: log residual = -33.263328, log improvement ratio=0.62  
6354  
[BiScaler] Iter 45: log residual = -33.889686, log improvement ratio=0.62  
6358  
[BiScaler] Iter 46: log residual = -34.516047, log improvement ratio=0.62  
6361  
[BiScaler] Iter 47: log residual = -35.142411, log improvement ratio=0.62  
6364  
[BiScaler] Iter 48: log residual = -35.768777, log improvement ratio=0.62  
6366  
[BiScaler] Iter 49: log residual = -36.395144, log improvement ratio=0.62  
6367  
[BiScaler] Iter 50: log residual = -37.021512, log improvement ratio=0.62  
6368  
[BiScaler] Iter 51: log residual = -37.647882, log improvement ratio=0.62  
6369  
[BiScaler] Iter 52: log residual = -38.274252, log improvement ratio=0.62  
6370  
[BiScaler] Iter 53: log residual = -38.900623, log improvement ratio=0.62  
6370  
[BiScaler] Iter 54: log residual = -39.526994, log improvement ratio=0.62  
6371  
[BiScaler] Iter 55: log residual = -40.153366, log improvement ratio=0.62  
6372

[BiScaler] Iter 56: log residual = -40.779738, log improvement ratio=0.62  
6372  
[BiScaler] Iter 57: log residual = -41.406109, log improvement ratio=0.62  
6372  
[BiScaler] Iter 58: log residual = -42.032481, log improvement ratio=0.62  
6372  
[BiScaler] Iter 59: log residual = -42.658853, log improvement ratio=0.62  
6372  
[BiScaler] Iter 60: log residual = -43.285226, log improvement ratio=0.62  
6373  
[BiScaler] Iter 61: log residual = -43.911602, log improvement ratio=0.62  
6376  
[BiScaler] Iter 62: log residual = -44.537966, log improvement ratio=0.62  
6364  
[BiScaler] Iter 63: log residual = -45.164347, log improvement ratio=0.62  
6381  
[BiScaler] Iter 64: log residual = -45.790719, log improvement ratio=0.62  
6373  
[BiScaler] Iter 65: log residual = -46.417087, log improvement ratio=0.62  
6368  
[BiScaler] Iter 66: log residual = -47.043472, log improvement ratio=0.62  
6385  
[BiScaler] Iter 67: log residual = -47.669846, log improvement ratio=0.62  
6373  
[BiScaler] Iter 68: log residual = -48.296203, log improvement ratio=0.62  
6357  
[BiScaler] Iter 69: log residual = -48.922611, log improvement ratio=0.62  
6409  
[BiScaler] Iter 70: log residual = -49.548918, log improvement ratio=0.62  
6307  
[BiScaler] Iter 71: log residual = -50.175297, log improvement ratio=0.62  
6378  
[BiScaler] Iter 72: log residual = -50.801673, log improvement ratio=0.62  
6376  
[BiScaler] Iter 73: log residual = -51.428247, log improvement ratio=0.62  
6575  
[BiScaler] Iter 74: log residual = -52.054331, log improvement ratio=0.62  
6084  
[BiScaler] Iter 75: log residual = -52.680838, log improvement ratio=0.62  
6507  
[BiScaler] Iter 76: log residual = -53.306945, log improvement ratio=0.62  
6106  
[BiScaler] Iter 77: log residual = -53.933222, log improvement ratio=0.62  
6278  
[BiScaler] Iter 78: log residual = -54.559726, log improvement ratio=0.62  
6504  
[BiScaler] Iter 79: log residual = -55.187109, log improvement ratio=0.62  
7383  
[BiScaler] Iter 80: log residual = -55.813326, log improvement ratio=0.62  
6217  
[BiScaler] Iter 81: log residual = -56.439541, log improvement ratio=0.62  
6215  
[BiScaler] Iter 82: log residual = -57.064246, log improvement ratio=0.62  
4705  
[BiScaler] Iter 83: log residual = -57.694122, log improvement ratio=0.62  
9876  
[BiScaler] Iter 84: log residual = -58.320642, log improvement ratio=0.62

```
6519
[BiScaler] Iter 85: log residual = -58.942397, log improvement ratio=0.62
1756
[BiScaler] Iter 86: log residual = -59.567736, log improvement ratio=0.62
5339
[BiScaler] Iter 87: log residual = -60.192510, log improvement ratio=0.62
4774
[BiScaler] Iter 88: log residual = -60.820491, log improvement ratio=0.62
7981
[BiScaler] Iter 89: log residual = -61.426252, log improvement ratio=0.60
5760
[BiScaler] Iter 90: log residual = -62.053421, log improvement ratio=0.62
7169
[BiScaler] Iter 91: log residual = -62.688779, log improvement ratio=0.63
5359
[BiScaler] Iter 92: log residual = -63.261398, log improvement ratio=0.57
2619
[BiScaler] Iter 93: log residual = -63.846200, log improvement ratio=0.58
4802
[BiScaler] Iter 94: log residual = -64.449178, log improvement ratio=0.60
2978
[BiScaler] Iter 95: log residual = -64.830932, log improvement ratio=0.38
1755
[BiScaler] Iter 96: log residual = -65.351248, log improvement ratio=0.52
0316
[BiScaler] Iter 97: log residual = -65.519441, log improvement ratio=0.16
8192
[BiScaler] Iter 98: log residual = -65.800609, log improvement ratio=0.28
1169
[BiScaler] Iter 99: log residual = -65.967886, log improvement ratio=0.16
7276
[BiScaler] Iter 100: log residual = -66.089223, log improvement ratio=0.1
21337
```

```
In [14]: # Access the alpha (row_mean) and beta (column) values
alpha = song_biscaler.row_means
alpha
```

```
Out[14]: array([ 0.23597732,  0.37836219, -0.12202861, ..., -0.00099498,
                -0.09292935, -0.07297946])
```

```
In [15]: beta = song_biscaler.column_means
beta
```

```
Out[15]: array([1.1085277 , 1.28189302, 1.45665654, 1.15972853, 1.28211814,
1.12983062, 1.47101127, 1.28577961, 1.35033994, 1.43749681,
1.16378442, 1.21511576, 1.4097627 , 1.23517437, 1.25240525,
1.18324648, 1.57097994, 1.24492046, 1.50054534, 1.21816032,
1.19678 , 1.45653583, 1.23245392, 1.1058456 , 1.41882857,
2.1370632 , 1.40920159, 1.12776456, 1.32721814, 1.21159835,
1.11964289, 1.20034543, 1.17693884, 1.58421436, 1.45240405,
1.27537058, 1.23174929, 1.15390764, 1.29276939, 1.11644992,
1.37531285, 1.14228572, 1.21817148, 1.26869206, 1.20196407,
1.18522847, 1.2132076 , 1.20720276, 1.12599004, 1.17277644,
1.26637501, 1.19294522, 1.16589595, 2.15784144, 1.7676486 ,
1.19068097, 1.12156866, 1.13149055, 1.85976324, 1.17344651,
1.41308493, 1.10785718, 1.09833489, 1.35995412, 1.31486195,
1.2667538 , 1.37337888, 1.11536594, 1.24306441, 1.45978565,
1.16395292, 1.21780274, 1.26713043, 1.20356712, 1.20973988,
1.4430925 , 1.23889642, 1.3110657 , 1.39755168, 1.35581768,
1.15864186, 1.19324072, 1.13430115, 1.14487043, 1.30399132,
1.13730439, 1.15994019, 1.33331344, 1.27431343, 1.31721697,
1.28752372, 1.30361275, 1.58147094, 1.24530107, 1.2061634 ,
1.2710075 , 1.27755075, 1.12600000, 1.11510750, 1.26056105])
```

```
In [16]: train.sort_index()
train_n = train.pivot_table(index="songID")
train_n['beta'] = beta
sorted = train_n.sort_values('beta', ascending=False)
sorted
```

```
Out[16]:
```

|        | rating   | userID      | year | beta     |
|--------|----------|-------------|------|----------|
| songID |          |             |      |          |
| 54     | 2.157240 | 1180.916137 | 1990 | 2.157841 |
| 26     | 2.134060 | 1196.513648 | 2001 | 2.137063 |
| 439    | 2.016957 | 1222.037037 | 2009 | 2.038437 |
| 637    | 1.950094 | 1221.344209 | 2002 | 1.937541 |
| 600    | 1.877886 | 1182.569444 | 2008 | 1.889610 |
| ...    | ...      | ...         | ...  | ...      |
| 317    | 1.077181 | 1119.037879 | 1995 | 1.073712 |
| 157    | 1.144404 | 1181.421769 | 2009 | 1.073472 |
| 165    | 1.069733 | 1316.726619 | 1996 | 1.072993 |
| 96     | 1.197060 | 1242.075188 | 2005 | 1.071808 |
| 481    | 1.132071 | 1186.690476 | 1978 | 1.062140 |

807 rows × 4 columns

```
In [17]: train[train['songID'].isin([54, 26, 439])]
```

```
Out[17]:
```

|        | userID | songID | rating   | songName       | year | artist        | genre   |
|--------|--------|--------|----------|----------------|------|---------------|---------|
| 11     | 2352   | 439    | 1.784426 | Secrets        | 2009 | OneRepublic   | Rock    |
| 229    | 1025   | 26     | 1.494918 | Undo           | 2001 | Bjork         | Rock    |
| 231    | 695    | 54     | 2.149164 | You're The One | 1990 | Dwight Yoakam | Country |
| 268    | 357    | 54     | 3.022958 | You're The One | 1990 | Dwight Yoakam | Country |
| 762    | 1284   | 439    | 1.784426 | Secrets        | 2009 | OneRepublic   | Rock    |
| ...    | ...    | ...    | ...      | ...            | ...  | ...           | ...     |
| 245482 | 1034   | 26     | 1.000000 | Undo           | 2001 | Bjork         | Rock    |
| 245558 | 650    | 54     | 1.494918 | You're The One | 1990 | Dwight Yoakam | Country |
| 245731 | 1873   | 439    | 1.000000 | Secrets        | 2009 | OneRepublic   | Rock    |
| 245739 | 1605   | 54     | 2.149164 | You're The One | 1990 | Dwight Yoakam | Country |
| 245790 | 833    | 26     | 4.073169 | Undo           | 2001 | Bjork         | Rock    |

2619 rows × 7 columns

```
songID songName Year Artist Genre Beta
54 You're The One 1990 Dwight Yoakam Country 2.157841
26 Undo 2001 Bjork Rock 2.137063
439 Secrets 2009 OneRepublic Rock 2.038437
```

The three largest beta's are the three most popular songs after removing for the bias due to how particular users rate certain songs. So here the 3 largest Bj values equal the 3 most popular songs

**b) iii.**



```
In [18]: train.sort_index()
train_a = train.pivot_table(index="userID")
train_a['alpha'] = alpha
# train_a
sorted_a = train_a.sort_values('alpha', ascending=False)
sorted_a
```

Out[18]:

|        | rating   | songID     | year        | alpha     |
|--------|----------|------------|-------------|-----------|
| userID |          |            |             |           |
| 1540   | 2.353383 | 395.373494 | 2003.945783 | 0.967014  |
| 838    | 2.105060 | 398.661157 | 2005.669421 | 0.814727  |
| 1569   | 2.077348 | 391.358491 | 2005.858491 | 0.794237  |
| 950    | 2.112970 | 426.039604 | 2003.782178 | 0.691602  |
| 345    | 2.109287 | 420.044248 | 2003.017699 | 0.674504  |
| ...    | ...      | ...        | ...         | ...       |
| 2076   | 1.059971 | 426.614035 | 2004.421053 | -0.404973 |
| 1550   | 1.063416 | 373.462687 | 2003.925373 | -0.405059 |
| 1595   | 1.014345 | 403.652174 | 2004.086957 | -0.408122 |
| 241    | 1.006971 | 390.014085 | 2005.450704 | -0.409727 |
| 2413   | 1.079893 | 413.809524 | 2003.174603 | -0.431814 |

2421 rows × 4 columns

```
userID alpha
1540 0.967014
838 0.814727
1569 0.794237
```

## b) iv.

```
In [19]: user_df = train_df.reset_index()[['userID']]
```

```
In [20]: test.sort_index()
test_df = test.pivot_table(index="userID", columns = "songID", values = "ra
```

```
In [21]: test_df=pd.merge(user_df,test_df.reset_index(),how='outer',on='userID')
test_df=test_df.set_index('userID')
test_mat = test_df.to_numpy()
test_mat
```

```
Out[21]: array([[nan, nan, nan, ..., nan, nan, nan],
               [nan, nan, nan, ..., nan, nan, nan],
               [nan, nan, nan, ..., nan, nan, nan],
               ...,
               [nan, nan, nan, ..., nan, 1., nan],
               [nan, nan, nan, ..., nan, nan, nan],
               [nan, nan, nan, ..., nan, nan, nan]])
```

```
In [22]: test_mask = ~np.isnan(test_mat)
print(np.sum(test_mask))
```

14471

```
In [23]: import copy
song_centered_0 = copy.copy(train_mat_centered)
song_centered_0[np.isnan(song_centered_0)]=0
song_centered_0
```

```
Out[23]: array([[0.          , 0.          , 0.          , ..., 0.          , 0.          ,
                0.          ],
               [0.          , 0.32958053, 0.          , ..., 0.          , 0.          ,
                0.          ],
               [0.          , 0.          , 0.          , ..., 0.          , 0.          ,
                0.          ],
               ...,
               [0.          , 0.          , 0.          , ..., 0.          , 0.          ,
                0.          ],
               [0.          , 0.          , 0.          , ..., 0.          , 0.          ,
                0.          ],
               [0.          , 0.28600431, 0.11124078, ..., 0.          , 0.          ,
                0.          ]])
```

```
In [24]: song_BiScaler_filled = song_biscaler.inverse_transform(song_centered_0)
song_BiScaler_filled
```

```
Out[24]: array([[1.34450503, 1.51787035, 1.69263387, ..., 1.55709616, 1.41842513,
                1.40129053],
               [1.4868899 , 1.98983574, 1.83501874, ..., 1.69948103, 1.56081   ,
                1.5436754 ],
               [0.98649909, 1.15986441, 1.33462793, ..., 1.19909023, 1.0604192 ,
                1.0432846 ],
               ...,
               [1.10753272, 1.28089804, 1.45566156, ..., 1.32012386, 1.18145283,
                1.16431823],
               [1.01559836, 1.18896367, 1.3637272 , ..., 1.22818949, 1.08951846,
                1.07238386],
               [1.03554825, 1.49491787, 1.49491787, ..., 1.24813938, 1.10946835,
                1.09233375]])
```

```
In [25]: def masked_mae(X_true, X_pred, mask):
        masked_diff = X_true[mask] - X_pred[mask]
        return np.mean(np.abs(masked_diff))

def masked_mse(X_true, X_pred, mask):
    masked_diff = X_true[mask] - X_pred[mask]
    return np.mean(masked_diff ** 2)

def OSR2(mse_model, mse_baseline):
    return 1 - mse_model/mse_baseline
```

```
In [26]: test_mae = masked_mae(test_mat, song_BiScaler_filled, test_mask)
print("Biscale MAE %s " % (test_mae/4)) #Note that we normalize MAE and RMS

test_mse = masked_mse(test_mat, song_BiScaler_filled, test_mask)
print("Biscale RMSE %s " % (np.sqrt(test_mse)/4))

baseline_pred = np.mean(train)[2]
baseline_model = baseline_pred*np.ones((2421, 807))
baseline_mse = masked_mse(test_mat, baseline_model, test_mask)
print("Biscale OSR2 %s" % OSR2(test_mse, baseline_mse))

Biscale MAE 0.0747571775133017
Biscale RMSE 0.0985693901114982
Biscale OSR2 0.2702373901935097
```

## c) i.

There are a total of  $807 + 2421 + 807k + 2421k = 3228(k+1)$  parameters included in the model. We have to train the model with 245997 observations.

## c) ii.

```
In [27]: vala.sort_index()
vala_df = vala.pivot_table(index="userID", columns = "songID", values = "ra

In [28]: vala_df=pd.merge(user_df,vala_df.reset_index(),how='outer',on='userID')
vala_df=vala_df.set_index('userID')
vala_mat = vala_df.to_numpy()
vala_mat
```

```
Out[28]: array([[nan, nan, nan, ..., nan, nan, nan],
               [nan, nan, nan, ..., nan, nan, nan],
               [nan, nan, nan, ..., nan, nan, nan],
               ...,
               [nan, nan, nan, ..., nan, nan, nan],
               [nan, nan, nan, ..., nan, nan, nan],
               [nan, nan, nan, ..., nan, nan, nan]])
```

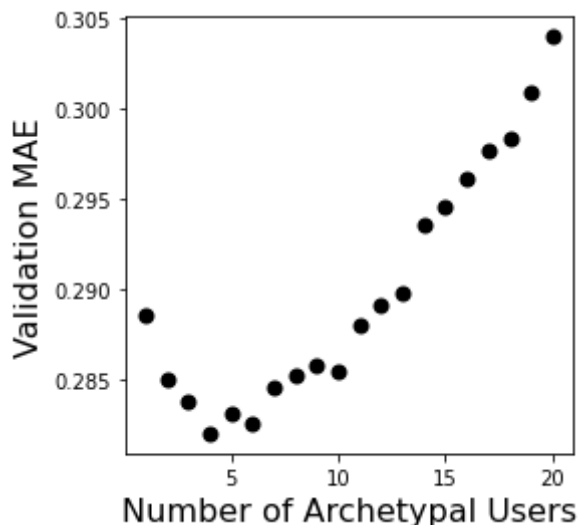
```
In [29]: vala_mask = ~np.isnan(vala_mat)
print(np.sum(vala_mask))
```

14470

```
In [30]: from fancyimpute import SoftImpute
songs_valA_mae_lst = []
for i in range(20):
    param_cv = i+1
    songs_soft_imputer_cv = SoftImpute(max_rank=param_cv, verbose=False)
    songs_centered_filled_cv = songs_soft_imputer_cv.fit_transform(train_ma
    songs_filled_cv = song_biscaler.inverse_transform(songs_centered_filled
    songs_filled_cv = np.clip(songs_filled_cv, 1, 5)
    songs_valA_mae_cv = masked_mae(vala_mat, songs_filled_cv, vala_mask)
    songs_valA_mae_lst.append(songs_valA_mae_cv)
    print('iter %s - Validation MAE %s' % (param_cv, songs_valA_mae_cv))
```

```
iter 1 - Validation MAE 0.2885634518667774
iter 2 - Validation MAE 0.28500121297797326
iter 3 - Validation MAE 0.2837867835747125
iter 4 - Validation MAE 0.2820231031496958
iter 5 - Validation MAE 0.2831439714205493
iter 6 - Validation MAE 0.2826094102603166
iter 7 - Validation MAE 0.28458407348601406
iter 8 - Validation MAE 0.2851844490181127
iter 9 - Validation MAE 0.2858040476209231
iter 10 - Validation MAE 0.28547984285183964
iter 11 - Validation MAE 0.2879752234507606
iter 12 - Validation MAE 0.28910102265746274
iter 13 - Validation MAE 0.2898415855585168
iter 14 - Validation MAE 0.2935896748594369
iter 15 - Validation MAE 0.2945718362743434
iter 16 - Validation MAE 0.29618740759125856
iter 17 - Validation MAE 0.2976931208274731
iter 18 - Validation MAE 0.2983513563550037
iter 19 - Validation MAE 0.30094743761703024
iter 20 - Validation MAE 0.30400269235062105
```

```
In [31]: import matplotlib.pyplot as plt
x = range(1,21)
y = songs_valA_mae_lst
plt.figure(figsize=(4, 4))
plt.scatter(x, y, linewidth=2, color='black')
plt.xlabel('Number of Archetypal Users', fontsize=16)
plt.ylabel('Validation MAE', fontsize=16)
plt.show()
```



I trained the models with k from 1-20 and evaluated the performance on validation set A. I chose k=4 as the number of archetypes because it has the lowest Validation MAE on the validation set A.

### c) iii.

```
In [32]: songs_soft_imputer = SoftImpute(max_rank=4, verbose=False) #use the best
songs_centered_filled = songs_soft_imputer.fit_transform(train_mat_centered)
songs_filled_matrix = song_biscaler.inverse_transform(songs_centered_filled)
songs_filled_matrix = np.clip(songs_filled_matrix, 1, 5)
```

```
In [33]: test_mae = masked_mae(test_mat, songs_filled_matrix, test_mask)
print("MAE %s " % (test_mae/4)) #Note that we normalize MAE and RMSE by the

test_mse = masked_mse(test_mat, songs_filled_matrix, test_mask)
print("RMSE %s " % (np.sqrt(test_mse)/4))

baseline_pred = np.mean(train_df)[2]
baseline_model = baseline_pred*np.ones((2421, 807))
baseline_mse = masked_mse(test_mat, baseline_model, test_mask)
print("OSR2 %s" % OSR2(test_mse, baseline_mse))
```

```
MAE 0.07063690502316015
RMSE 0.09598278651787194
OSR2 0.30874774832022855
```

### d) i.

## Linear Regression

```
In [34]: train['genre'] = train.genre.astype('category')
train['year'] = train.year.astype('category')
test['genre'] = test.genre.astype('category')
test['year'] = test.year.astype('category')
test
```

Out[34]:

|       | userID | songID | rating   | songName              | year | artist                | genre   |
|-------|--------|--------|----------|-----------------------|------|-----------------------|---------|
| 0     | 853    | 54     | 1.989836 | You're The One        | 1990 | Dwight Yoakam         | Country |
| 1     | 608    | 34     | 1.000000 | Nah!                  | 2002 | Shania Twain          | Country |
| 2     | 9      | 54     | 3.969507 | You're The One        | 1990 | Dwight Yoakam         | Country |
| 3     | 1862   | 54     | 2.279344 | You're The One        | 1990 | Dwight Yoakam         | Country |
| 4     | 329    | 80     | 1.494918 | Tim McGraw            | 2006 | Taylor Swift          | Country |
| ...   | ...    | ...    | ...      | ...                   | ...  | ...                   | ...     |
| 14466 | 443    | 615    | 1.494918 | Daylight              | 2002 | Coldplay              | Rock    |
| 14467 | 1061   | 657    | 1.000000 | Kissy Kissy           | 2003 | The Kills             | Rock    |
| 14468 | 610    | 653    | 1.000000 | The Laws Have Changed | 2003 | The New Pornographers | Rock    |
| 14469 | 690    | 657    | 1.000000 | Kissy Kissy           | 2003 | The Kills             | Rock    |
| 14470 | 230    | 657    | 1.000000 | Kissy Kissy           | 2003 | The Kills             | Rock    |

14471 rows × 7 columns

```
In [35]: import statsmodels.formula.api as smf
ols = smf.ols(formula='rating ~ genre+year', data=train)
res = ols.fit()
print(res.summary())
```

```

                                OLS Regression Results
=====
Dep. Variable:                  rating    R-squared:
0.034
Model:                          OLS      Adj. R-squared:
0.034
Method:                        Least Squares    F-statistic:
261.6
Date:                          Wed, 28 Apr 2021    Prob (F-statistic):
0.00
Time:                          13:56:00    Log-Likelihood:
3e+05                                -1.537
No. Observations:              245997    AIC:
5e+05                                3.07
Df Residuals:                  245963    BIC:
9e+05                                3.07
Df Model:                      33
Covariance Type:               nonrobust
=====
=====
                                coef    std err          t      P>|t|      [0.0
25      0.975]
-----
-----
Intercept                2.0937      0.024     87.425      0.000      2.0
47      2.141
genre[T.Electronic]     -0.0777      0.011    -6.813      0.000     -0.1
00      -0.055
genre[T.Folk]           -0.0974      0.012    -7.881      0.000     -0.1
22      -0.073
genre[T.Pop]            -0.1351      0.011   -12.056      0.000     -0.1
57      -0.113
genre[T.Rap]            -0.1118      0.012    -9.211      0.000     -0.1
36      -0.088
genre[T.RnB]            -0.2781      0.012   -23.772      0.000     -0.3
01      -0.255
genre[T.Rock]           -0.1704      0.011   -15.423      0.000     -0.1
92      -0.149
year[T.1976]            -0.4801      0.027   -17.484      0.000     -0.5
34      -0.426
year[T.1978]            -0.7912      0.041   -19.351      0.000     -0.8
71      -0.711
year[T.1979]            -0.1324      0.032    -4.110      0.000     -0.1
96      -0.069
year[T.1985]            -0.5300      0.043   -12.281      0.000     -0.6
15      -0.445
year[T.1986]            -0.4073      0.027   -15.162      0.000     -0.4
60      -0.355
year[T.1987]            -0.5412      0.028   -19.487      0.000     -0.5
96      -0.487
year[T.1988]            -0.3531      0.026   -13.483      0.000     -0.4

```

|              |        |         |       |         |       |      |
|--------------|--------|---------|-------|---------|-------|------|
| 04           | -0.302 |         |       |         |       |      |
| year[T.1990] |        | -0.0660 | 0.027 | -2.480  | 0.013 | -0.1 |
| 18           | -0.014 |         |       |         |       |      |
| year[T.1991] |        | -0.6384 | 0.025 | -25.474 | 0.000 | -0.6 |
| 88           | -0.589 |         |       |         |       |      |
| year[T.1992] |        | -0.5042 | 0.024 | -20.807 | 0.000 | -0.5 |
| 52           | -0.457 |         |       |         |       |      |
| year[T.1993] |        | -0.5394 | 0.024 | -22.354 | 0.000 | -0.5 |
| 87           | -0.492 |         |       |         |       |      |
| year[T.1995] |        | -0.4788 | 0.026 | -18.487 | 0.000 | -0.5 |
| 30           | -0.428 |         |       |         |       |      |
| year[T.1996] |        | -0.8069 | 0.023 | -35.725 | 0.000 | -0.8 |
| 51           | -0.763 |         |       |         |       |      |
| year[T.1997] |        | -0.5502 | 0.024 | -23.338 | 0.000 | -0.5 |
| 96           | -0.504 |         |       |         |       |      |
| year[T.1998] |        | -0.8162 | 0.027 | -30.583 | 0.000 | -0.8 |
| 69           | -0.764 |         |       |         |       |      |
| year[T.1999] |        | -0.5511 | 0.022 | -24.644 | 0.000 | -0.5 |
| 95           | -0.507 |         |       |         |       |      |
| year[T.2000] |        | -0.6126 | 0.022 | -28.162 | 0.000 | -0.6 |
| 55           | -0.570 |         |       |         |       |      |
| year[T.2001] |        | -0.5470 | 0.022 | -24.951 | 0.000 | -0.5 |
| 90           | -0.504 |         |       |         |       |      |
| year[T.2002] |        | -0.5946 | 0.022 | -27.436 | 0.000 | -0.6 |
| 37           | -0.552 |         |       |         |       |      |
| year[T.2003] |        | -0.6362 | 0.021 | -29.633 | 0.000 | -0.6 |
| 78           | -0.594 |         |       |         |       |      |
| year[T.2004] |        | -0.7157 | 0.022 | -32.678 | 0.000 | -0.7 |
| 59           | -0.673 |         |       |         |       |      |
| year[T.2005] |        | -0.6430 | 0.022 | -29.902 | 0.000 | -0.6 |
| 85           | -0.601 |         |       |         |       |      |
| year[T.2006] |        | -0.6379 | 0.021 | -29.753 | 0.000 | -0.6 |
| 80           | -0.596 |         |       |         |       |      |
| year[T.2007] |        | -0.6421 | 0.021 | -29.924 | 0.000 | -0.6 |
| 84           | -0.600 |         |       |         |       |      |
| year[T.2008] |        | -0.6450 | 0.021 | -30.112 | 0.000 | -0.6 |
| 87           | -0.603 |         |       |         |       |      |
| year[T.2009] |        | -0.5772 | 0.021 | -26.930 | 0.000 | -0.6 |
| 19           | -0.535 |         |       |         |       |      |
| year[T.2010] |        | -0.6543 | 0.022 | -30.022 | 0.000 | -0.6 |
| 97           | -0.612 |         |       |         |       |      |

=====

|                |           |                   |       |
|----------------|-----------|-------------------|-------|
| Omnibus:       | 63019.216 | Durbin-Watson:    |       |
| 2.003          |           |                   |       |
| Prob(Omnibus): | 0.000     | Jarque-Bera (JB): | 14508 |
| 5.217          |           |                   |       |
| Skew:          | 1.457     | Prob(JB):         |       |
| 0.00           |           |                   |       |
| Kurtosis:      | 5.380     | Cond. No.         |       |
| 153.           |           |                   |       |

=====

#### Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.



```
In [36]: test_pred_ols = res.predict(test)
test_mae_ols = np.mean(np.abs(test['rating'] - test_pred_ols))
print("MAE %s " % (test_mae_ols/4)) #Note that we normalize MAE and RMSE by

test_mse_ols = np.mean((test['rating'] - test_pred_ols)**2)
print("RMSE %s " % (np.sqrt(test_mse_ols)/4))

print("OSR2 %s " % OSR2(test_mse_ols, baseline_mse))

MAE 0.0923556815359249
RMSE 0.11312545968408343
OSR2 0.039780403459793834
```

Random Forest

```
In [41]: from sklearn.ensemble import RandomForestRegressor

X_train = train.drop(columns = ["genre", "songName", "artist"])
y_train = train["rating"]
rf = RandomForestRegressor(max_features=len(X_train.columns), min_samples_1
                           n_estimators = 500, random_state=88, verbose=2)
rf.fit(X_train, y_train)

[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent w
orkers.
[Parallel(n_jobs=1)]: Done    1 out of    1 | elapsed:    0.1s remaining:
0.0s
```

```
In [42]: test_pred_rf=rf.predict(test.drop(columns = ["genre","songName","artist"]))
test_mae_rf = np.mean(np.abs(test['rating'] - test_pred_rf))
print("MAE %s " % (test_mae_rf/4)) #Note that we normalize MAE and RMSE by

test_mse_rf = np.mean((test['rating'] - test_pred_rf)**2)
print("RMSE %s " % (np.sqrt(test_mse_rf)/4))

print("OSR2 %s " % OSR2(test_mse_rf, baseline_mse))

[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent w
orkers.
[Parallel(n_jobs=1)]: Done 1 out of 1 | elapsed: 0.0s remaining:
0.0s

MAE 3.0931830823804327e-06
RMSE 0.0002774117114856834
OSR2 0.9999942256987917

[Parallel(n_jobs=1)]: Done 500 out of 500 | elapsed: 0.3s finished
```

## d) ii.

```
In [ ]: val_pred_cf = train_filled_matrix[valB_mask]

blend_valB_df = songs_valB.reset_index()[['userID','movieID','rating']]

blend_valB_df['val_pred_cf']=val_pred_cf
blend_valB_df
```

```
In [ ]: val_pred_ols = res.predict(valB)
blend_valB_df['val_pred_ols']=val_pred_ols
blend_valB_df
```

```
In [ ]: blend_valB_df['val_pred_rf']=val_pred_rf
blend_valB_df
```

```
In [ ]: val_pred_blended =blending_res.predict(blend_valB_df)
```

```
In [ ]: test_mae_blended = np.mean(np.abs(test_mat[test_mask] - val_pred_blended))
print("Test_blended MAE %s " % (test_mae_blended/4)) #Note that we normali

test_mse_blended = np.mean((movie_lens_test_mat[movie_lens_test_mask] - tes
print("Test_blended RMSE %s " % (np.sqrt(test_mse_blended)/4))

print("Test_blended OSR2 %s " % OSR2(test_mse_blended, baseline_mse))
```